

A Practical Introduction to Control, Numerics and Machine Learning

Day 2

Summer School IFAC CPDE 2022

Workshop on Control of Systems Governed by Partial Differential Equations

Daniël Veldman

Chair in Dynamics, Control, and Numerics, Friedrich-Alexander-University Erlangen-Nürnberg

Contents

- 2.A** Deep (residual) neural networks and neural ODEs
- 2.B** Sensitivity analysis with neural ODEs
- 2.C** Training of deep residual neural networks



2.A Deep (residual) neural networks and neural ODEs




Neural network

Input layer

Hidden layer

Output layer



The number of nodes in the input layer is the number of inputs.

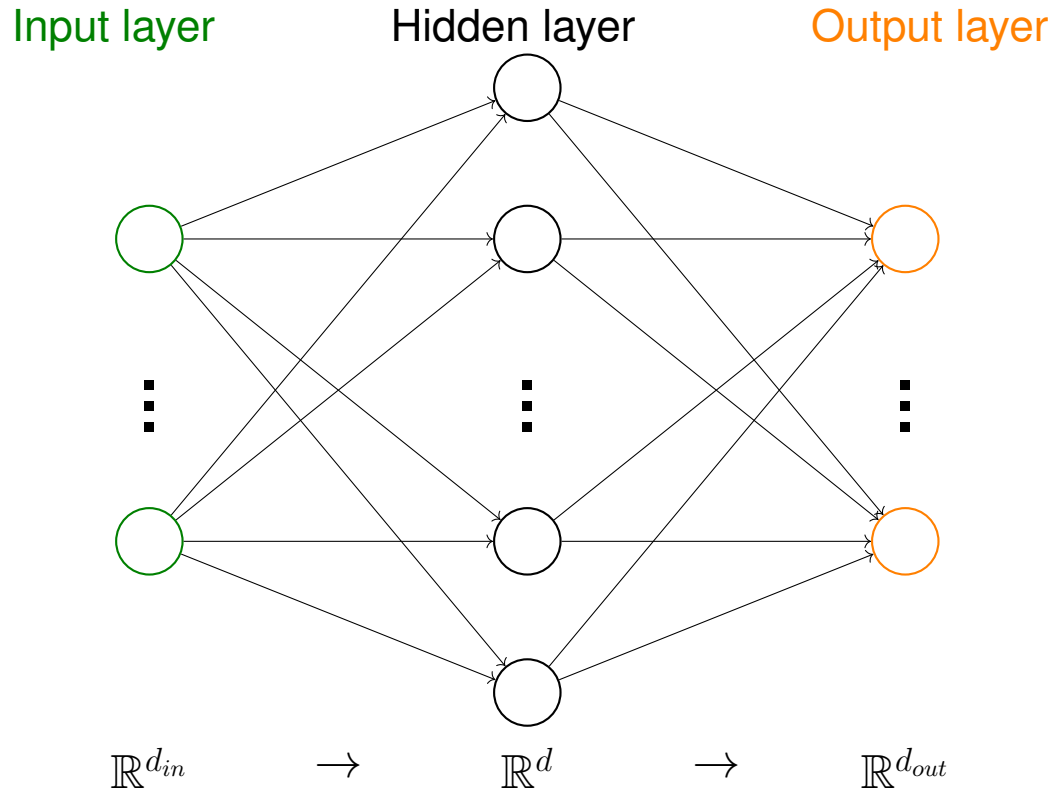
Each arrow represents a linear map.

Each node in the hidden layer represents a nonlinear operation $\sigma(\cdot)$.

The number of nodes in the output layer is the number of outputs.

$$\mathbf{y} = \mathbf{V}\sigma(\mathbf{W}\mathbf{x} + \mathbf{c}) + \mathbf{b}$$

Representation theorem



$$\mathbf{y} = \mathbf{V}\sigma(\mathbf{W}\mathbf{x} + \mathbf{c}) + \mathbf{b}$$

Any function $\mathbf{y} = f(\mathbf{x})$ (e.g. in L^2) can be approximated arbitrarily well by a sufficiently **wide** neural network $\mathbf{y} = \mathbf{V}\sigma(\mathbf{W}\mathbf{x} + \mathbf{c}) + \mathbf{b}$.

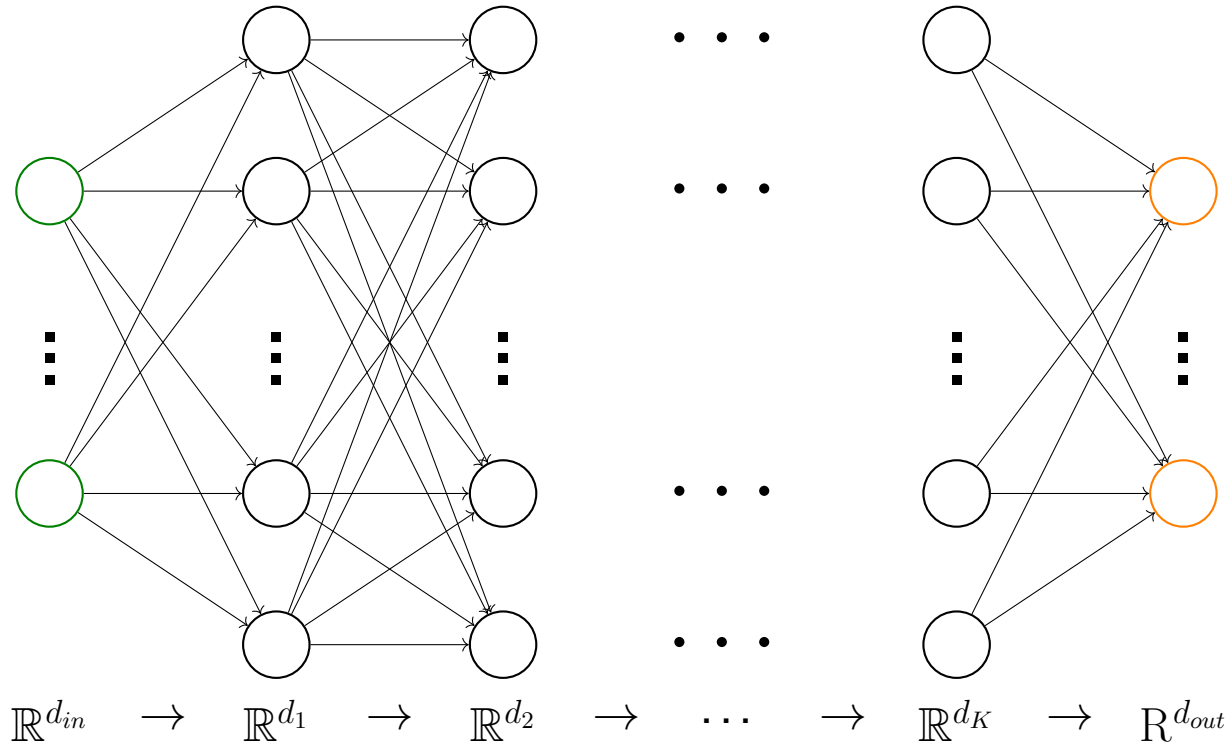
(Cybenko, 1989)

Deep Neural Networks

Input layer

K hidden layers

Output layer



There are now multiple hidden layers:

$$\mathbf{x}_0 = \mathbf{x}_{in}, \quad \mathbf{x}_k = \mathbf{V}_k \sigma(\mathbf{x}_{k-1}) + \mathbf{b}_k, \quad \mathbf{y} = \mathbf{x}_K.$$

Residual neural networks (ResNet) and neural ODEs

$$\mathbf{x}_0 = \mathbf{x}_{\text{in}}, \quad \mathbf{x}_k = \mathbf{V}_k \sigma(\mathbf{x}_{k-1}) + \mathbf{b}_k, \quad \mathbf{y} = \mathbf{x}_K.$$

When the number of hidden layers K is large,
it is better to consider a **residual neural network (ResNet)**

$$\mathbf{x}_0 = \mathbf{x}_{\text{in}}, \quad \mathbf{x}_k = \mathbf{x}_{k-1} + \mathbf{V}_k \sigma(\mathbf{x}_{k-1}) + \mathbf{b}_k, \quad \mathbf{y} = \mathbf{x}_K.$$

We can view a ResNet as Forward Euler discretization of the **neural ODE**:

$$\mathbf{x}(0) = \mathbf{x}_{\text{in}}, \quad \dot{\mathbf{x}}(t) = \mathbf{V}(t) \sigma(\mathbf{x}(t)) + \mathbf{b}(t), \quad \mathbf{y} = \mathbf{x}(T).$$

How do we find the weights $\mathbf{V}(t)$ and $\mathbf{b}(t)$?

Training a neural ODE

$$\mathbf{x}(0) = \mathbf{x}_{\text{in}}, \quad \dot{\mathbf{x}}(t) = \mathbf{V}(t)\sigma(\mathbf{x}(t)) + \mathbf{b}(t), \quad \mathbf{y} = \mathbf{x}(T).$$

Given training data: pairs $(\mathbf{x}_{\text{in}}^i, \mathbf{y}_{\text{out}}^i)$, $i = 1, 2, 3, \dots, I$.

$\mathbf{y}_{\text{out}}^i$ is the desired output for the input \mathbf{x}_{in}^i .

For certain weights $\mathbf{V}(t)$ and $\mathbf{b}(t)$, the output resulting from the input \mathbf{x}_{in}^i is $\mathbf{x}^i(T)$

We thus want to minimize

$$J(\mathbf{V}, \mathbf{b}) = \frac{1}{2} \sum_{i=1}^I |\mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i|^2 + \frac{w_1}{2} \sum_{i=1}^I \int_0^T |\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i|^2 dt + \frac{w_2}{2} \int_0^T (\|\mathbf{V}(t)\|_F^2 + |\mathbf{b}(t)|^2) dt.$$

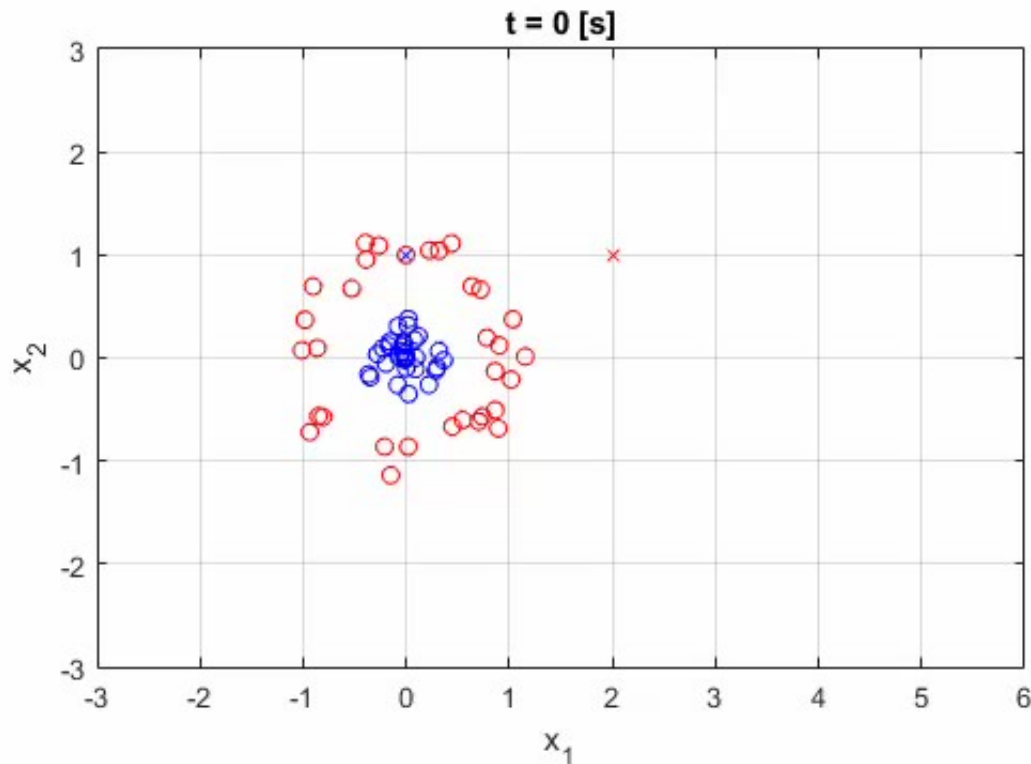
subject to the dynamics (for $i = 1, 2, 3, \dots, I$)

$$\mathbf{x}^i(0) = \mathbf{x}_{\text{in}}^i, \quad \dot{\mathbf{x}}^i(t) = \mathbf{V}(t)\sigma(\mathbf{x}^i(t)) + \mathbf{b}(t).$$

Note: $\mathbf{b}(t)$ is a vector, but $\mathbf{V}(t)$ is a (square) matrix.

Note: $\|\cdot\|_F$ denotes the Frobenius norm, $\|\mathbf{V}\|_F := \sqrt{\text{trace}(\mathbf{V}^\top \mathbf{V})}$.


An example: initial data x_{in}^i



blue data points
have target
 $y_{\text{out}}^i = (0, 1)$.

red data points
have target
 $y_{\text{out}}^i = (2, 1)$.


An example: evolution of $x^i(t)$



blue data points
have target
 $y_{\text{out}}^i = (0, 1)$.

red data points
have target
 $y_{\text{out}}^i = (2, 1)$.

An example: final data $x^i(T)$



blue data points
have target
 $y_{\text{out}}^i = (0, 1)$.

red data points
have target
 $y_{\text{out}}^i = (2, 1)$.

2.B Sensitivity analysis with neural ODEs



The directional derivative w.r.t. $\mathbf{b}(t)$

Consider a perturbation $\tilde{\mathbf{b}}(t)$ and compute:

$$\tilde{\mathbf{x}}^i(t) := \lim_{h \rightarrow 0} \frac{\mathbf{x}^{ih}(t) - \mathbf{x}^i(t)}{h}$$

where

$$\mathbf{x}^i(0) = \mathbf{x}_{\text{in}}^i,$$

$$\mathbf{x}^{ih}(0) = \mathbf{x}_{\text{in}}^i,$$

$$\dot{\mathbf{x}}^i(t) = \mathbf{V}(t)\sigma(\mathbf{x}^i(t)) + \mathbf{b}(t),$$

$$\dot{\mathbf{x}}^{ih}(t) = \mathbf{V}(t)\sigma(\mathbf{x}^{ih}(t)) + \mathbf{b}(t) + h\tilde{\mathbf{b}}(t).$$

The directional derivative w.r.t. $\mathbf{b}(t)$

Consider a perturbation $\tilde{\mathbf{b}}(t)$ and compute:

$$\tilde{\mathbf{x}}^i(t) := \lim_{h \rightarrow 0} \frac{\mathbf{x}^{ih}(t) - \mathbf{x}^i(t)}{h}$$

where

$$\mathbf{x}^i(0) = \mathbf{x}_{\text{in}}^i,$$

$$\mathbf{x}^{ih}(0) = \mathbf{x}_{\text{in}}^i,$$

$$\dot{\mathbf{x}}^i(t) = \mathbf{V}(t)\sigma(\mathbf{x}^i(t)) + \mathbf{b}(t),$$

$$\dot{\mathbf{x}}^{ih}(t) = \mathbf{V}(t)\sigma(\mathbf{x}^{ih}(t)) + \mathbf{b}(t) + h\tilde{\mathbf{b}}(t).$$

Because

$$\frac{\dot{\mathbf{x}}^{ih} - \dot{\mathbf{x}}^i}{h} = \frac{\mathbf{V}\sigma(\mathbf{x}^{ih}) + \mathbf{b} + h\tilde{\mathbf{b}} - \mathbf{V}\sigma(\mathbf{x}^i) - \mathbf{b}}{h} = \mathbf{V} \frac{\sigma(\mathbf{x}^{ih}) - \sigma(\mathbf{x}^i)}{h} + \tilde{\mathbf{b}},$$

taking the limit $h \rightarrow 0$ shows that $\tilde{\mathbf{x}}^i(t)$ satisfies

$$\tilde{\mathbf{x}}^i(0) = 0, \quad \dot{\tilde{\mathbf{x}}}^i(t) = \mathbf{V}(t) \text{diag} \left(\frac{d\sigma}{dx}(\mathbf{x}^i(t)) \right) \tilde{\mathbf{x}}^i(t) + \tilde{\mathbf{b}}(t).$$

The directional derivative w.r.t. $\mathbf{b}(t)$

Consider a perturbation $\tilde{\mathbf{b}}(t)$ and compute:

$$\tilde{\mathbf{x}}^i(t) := \lim_{h \rightarrow 0} \frac{\mathbf{x}^{ih}(t) - \mathbf{x}^i(t)}{h}$$

where

$$\begin{aligned} \mathbf{x}^i(0) &= \mathbf{x}_{\text{in}}^i, & \dot{\mathbf{x}}^i(t) &= \mathbf{V}(t)\sigma(\mathbf{x}^i(t)) + \mathbf{b}(t), \\ \mathbf{x}^{ih}(0) &= \mathbf{x}_{\text{in}}^i, & \dot{\mathbf{x}}^{ih}(t) &= \mathbf{V}(t)\sigma(\mathbf{x}^{ih}(t)) + \mathbf{b}(t) + h\tilde{\mathbf{b}}(t). \end{aligned}$$

Because

$$\frac{\dot{\mathbf{x}}^{ih} - \dot{\mathbf{x}}^i}{h} = \frac{\mathbf{V}\sigma(\mathbf{x}^{ih}) + \mathbf{b} + h\tilde{\mathbf{b}} - \mathbf{V}\sigma(\mathbf{x}^i) - \mathbf{b}}{h} = \mathbf{V} \frac{\sigma(\mathbf{x}^{ih}) - \sigma(\mathbf{x}^i)}{h} + \tilde{\mathbf{b}},$$

taking the limit $h \rightarrow 0$ shows that $\tilde{\mathbf{x}}^i(t)$ satisfies

$$\tilde{\mathbf{x}}^i(0) = 0, \quad \dot{\tilde{\mathbf{x}}}^i(t) = \mathbf{V}(t) \text{diag} \left(\frac{d\sigma}{dx}(\mathbf{x}^i(t)) \right) \tilde{\mathbf{x}}^i(t) + \tilde{\mathbf{b}}(t).$$

We then obtain

$$\begin{aligned} \langle \nabla_{\mathbf{b}} J, \tilde{\mathbf{b}} \rangle &= \lim_{h \rightarrow 0} \frac{J(\mathbf{V}, \mathbf{b} + h\tilde{\mathbf{b}}) - J(\mathbf{V}, \mathbf{b})}{h} = \sum_{i=1}^I (\mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i)^\top \tilde{\mathbf{x}}^i(T) + \\ &\quad w_1 \sum_{i=1}^I \int_0^T (\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i)^\top \tilde{\mathbf{x}}^i(t) dt + w_2 \int_0^T (\mathbf{b}(t))^\top \tilde{\mathbf{b}}(t) dt. \end{aligned}$$

The directional derivative w.r.t. $\mathbf{V}(t)$

Consider a perturbation $\hat{\mathbf{V}}(t)$ and compute:

$$\hat{\mathbf{x}}^i(t) := \lim_{h \rightarrow 0} \frac{\mathbf{x}^{ih}(t) - \mathbf{x}^i(t)}{h}$$

where

$$\mathbf{x}^i(0) = \mathbf{x}_{\text{in}}^i,$$

$$\mathbf{x}^{ih}(0) = \mathbf{x}_{\text{in}}^i,$$

$$\dot{\mathbf{x}}^i(t) = \mathbf{V}(t)\sigma(\mathbf{x}^i(t)) + \mathbf{b}(t),$$

$$\dot{\mathbf{x}}^{ih}(t) = (\mathbf{V}(t) + h\hat{\mathbf{V}}(t))\sigma(\mathbf{x}^{ih}(t)) + \mathbf{b}(t).$$

The directional derivative w.r.t. $\mathbf{V}(t)$

Consider a perturbation $\hat{\mathbf{V}}(t)$ and compute:

$$\hat{\mathbf{x}}^i(t) := \lim_{h \rightarrow 0} \frac{\mathbf{x}^{ih}(t) - \mathbf{x}^i(t)}{h}$$

where

$$\mathbf{x}^i(0) = \mathbf{x}_{\text{in}}^i,$$

$$\mathbf{x}^{ih}(0) = \mathbf{x}_{\text{in}}^i,$$

$$\dot{\mathbf{x}}^i(t) = \mathbf{V}(t)\sigma(\mathbf{x}^i(t)) + \mathbf{b}(t),$$

$$\dot{\mathbf{x}}^{ih}(t) = (\mathbf{V}(t) + h\hat{\mathbf{V}}(t))\sigma(\mathbf{x}^{ih}(t)) + \mathbf{b}(t).$$

Because

$$\frac{\dot{\mathbf{x}}^{ih} - \dot{\mathbf{x}}^i}{h} = \frac{(\mathbf{V} + h\hat{\mathbf{V}})\sigma(\mathbf{x}^{ih}) + \mathbf{b} - \mathbf{V}\sigma(\mathbf{x}^{ih}) + \mathbf{V}\sigma(\mathbf{x}^{ih}) - \mathbf{V}\sigma(\mathbf{x}^i) - \mathbf{b}}{h},$$

taking the limit $h \rightarrow 0$ shows that $\hat{\mathbf{x}}^i(t)$ satisfies

$$\hat{\mathbf{x}}^i(0) = 0, \quad \dot{\hat{\mathbf{x}}}^i(t) = \hat{\mathbf{V}}(t)\sigma(\mathbf{x}^i(t)) + \mathbf{V}(t)\text{diag}\left(\frac{d\sigma}{dx}(\mathbf{x}^i(t))\right)\hat{\mathbf{x}}^i(t).$$

The directional derivative w.r.t. $\mathbf{V}(t)$

Consider a perturbation $\hat{\mathbf{V}}(t)$ and compute:

$$\hat{\mathbf{x}}^i(t) := \lim_{h \rightarrow 0} \frac{\mathbf{x}^{ih}(t) - \mathbf{x}^i(t)}{h}$$

where

$$\mathbf{x}^i(0) = \mathbf{x}_{\text{in}}^i,$$

$$\dot{\mathbf{x}}^i(t) = \mathbf{V}(t)\sigma(\mathbf{x}^i(t)) + \mathbf{b}(t),$$

$$\mathbf{x}^{ih}(0) = \mathbf{x}_{\text{in}}^i,$$

$$\dot{\mathbf{x}}^{ih}(t) = (\mathbf{V}(t) + h\hat{\mathbf{V}}(t))\sigma(\mathbf{x}^{ih}(t)) + \mathbf{b}(t).$$

Because

$$\frac{\dot{\mathbf{x}}^{ih} - \dot{\mathbf{x}}^i}{h} = \frac{(\mathbf{V} + h\hat{\mathbf{V}})\sigma(\mathbf{x}^{ih}) + \mathbf{b} - \mathbf{V}\sigma(\mathbf{x}^{ih}) + \mathbf{V}\sigma(\mathbf{x}^{ih}) - \mathbf{V}\sigma(\mathbf{x}^i) - \mathbf{b}}{h},$$

taking the limit $h \rightarrow 0$ shows that $\hat{\mathbf{x}}^i(t)$ satisfies

$$\hat{\mathbf{x}}^i(0) = 0, \quad \dot{\hat{\mathbf{x}}}^i(t) = \hat{\mathbf{V}}(t)\sigma(\mathbf{x}^i(t)) + \mathbf{V}(t)\text{diag}\left(\frac{d\sigma}{dx}(\mathbf{x}^i(t))\right) \hat{\mathbf{x}}^i(t).$$

We then obtain

$$\begin{aligned} \langle \nabla_{\mathbf{V}} J, \hat{\mathbf{V}} \rangle_F &= \lim_{h \rightarrow 0} \frac{J(\mathbf{V} + h\hat{\mathbf{V}}, \mathbf{b}) - J(\mathbf{V}, \mathbf{b})}{h} = \sum_{i=1}^I (\mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i)^\top \hat{\mathbf{x}}^i(T) + \\ &w_1 \sum_{i=1}^I \int_0^T (\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i)^\top \hat{\mathbf{x}}^i(t) dt + w_2 \int_0^T \langle \mathbf{V}(t), \hat{\mathbf{V}}(t) \rangle_F dt. \end{aligned}$$

The adjoint state

Just as in the previous lecture, we need the adjoint state to find the gradients.

We now define the adjoint states $\varphi^i(t)$:

$$\varphi^i(T) = \mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i, \quad -\dot{\varphi}^i(t) = \left(\mathbf{V}(t) \text{diag} \left(\frac{d\sigma}{dx}(\mathbf{x}^i(t)) \right) \right)^\top \varphi^i(t) + w_1(\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i),$$

for $i = 1, 2, 3, \dots, I$.

Note: the final condition is now nonzero because the state $\mathbf{x}^i(T)$ at the final time appears in the cost functional.

Question 1

Write:

$$\mathbf{A}(t) = \mathbf{V}(t) \text{diag} \left(\frac{d\sigma}{dx}(\mathbf{x}^i(t)) \right).$$

From the previous slides

$$\tilde{\mathbf{x}}^i(0) = 0, \quad \dot{\tilde{\mathbf{x}}}^i(t) = \mathbf{A}(t)\tilde{\mathbf{x}}^i(t) + \tilde{\mathbf{b}}(t).$$

$$\boldsymbol{\varphi}^i(T) = \mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i, \quad -\dot{\boldsymbol{\varphi}}^i(t) = (\mathbf{A}(t))^\top \boldsymbol{\varphi}^i(t) + w_1(\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i).$$

What is

$$\int_0^T \frac{d}{dt} \left((\boldsymbol{\varphi}^i(t))^\top \tilde{\mathbf{x}}^i(t) \right) dt = (\boldsymbol{\varphi}^i(T))^\top \tilde{\mathbf{x}}^i(T) - (\boldsymbol{\varphi}^i(0))^\top \tilde{\mathbf{x}}^i(0)?$$

A) 0

B) $\mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i$

C) $(\mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i)^\top \tilde{\mathbf{x}}^i(T)$

D) $-(\mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i)^\top \tilde{\mathbf{x}}^i(T)$

E) None of the above

Question 2

We have that:

$$\tilde{\mathbf{x}}^i(0) = 0, \quad \dot{\tilde{\mathbf{x}}}^i(t) = \mathbf{A}(t)\tilde{\mathbf{x}}^i(t) + \tilde{\mathbf{b}}(t).$$

$$\boldsymbol{\varphi}^i(T) = \mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i, \quad -\dot{\boldsymbol{\varphi}}^i(t) = (\mathbf{A}(t))^\top \boldsymbol{\varphi}^i(t) + w_1(\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i).$$

What is

$$\int_0^T \frac{d}{dt} \left((\boldsymbol{\varphi}^i(t))^\top \tilde{\mathbf{x}}^i(t) \right) dt = \int_0^T (\dot{\boldsymbol{\varphi}}^i(t))^\top \tilde{\mathbf{x}}^i(t) dt + \int_0^T (\boldsymbol{\varphi}^i(t))^\top \dot{\tilde{\mathbf{x}}}^i(t) dt?$$

A) $\int_0^T \left((\mathbf{A}(t))^\top \boldsymbol{\varphi}^i(t) + w_1(\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i) \right)^\top \tilde{\mathbf{x}}^i(t) dt + \int_0^T (\boldsymbol{\varphi}^i(t))^\top \left(\mathbf{A}(t)\tilde{\mathbf{x}}^i(t) + \tilde{\mathbf{b}}(t) \right) dt$

B) $-\int_0^T \left((\mathbf{A}(t))^\top \boldsymbol{\varphi}^i(t) + w_1(\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i) \right)^\top \tilde{\mathbf{x}}^i(t) dt + \int_0^T (\boldsymbol{\varphi}^i(t))^\top \left(\mathbf{A}(t)\tilde{\mathbf{x}}^i(t) + \tilde{\mathbf{b}}(t) \right) dt$

C) $w_1 \int_0^T (\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i)^\top \tilde{\mathbf{x}}^i(t) dt + \int_0^T (\boldsymbol{\varphi}^i(t))^\top \tilde{\mathbf{b}}(t) dt$

D) $-w_1 \int_0^T (\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i)^\top \tilde{\mathbf{x}}^i(t) dt + \int_0^T (\boldsymbol{\varphi}^i(t))^\top \tilde{\mathbf{b}}(t) dt$

The gradient w.r.t. $\mathbf{b}(t)$

$$\langle \nabla_{\mathbf{b}} J, \tilde{\mathbf{b}} \rangle = \lim_{h \rightarrow 0} \frac{J(\mathbf{V}, \mathbf{b} + h\tilde{\mathbf{b}}) - J(\mathbf{V}, \mathbf{b})}{h} = \sum_{i=1}^I (\mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i)^\top \tilde{\mathbf{x}}^i(T) +$$

$$w_1 \sum_{i=1}^I \int_0^T (\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i)^\top \tilde{\mathbf{x}}^i(t) dt + w_2 \int_0^T (\mathbf{b}(t))^\top \tilde{\mathbf{b}}(t) dt.$$

From the previous questions:

$$(\mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i)^\top \tilde{\mathbf{x}}^i(T) = -w_1 \int_0^T (\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i)^\top \tilde{\mathbf{x}}^i(t) dt + \int_0^T (\boldsymbol{\varphi}^i(t))^\top \tilde{\mathbf{b}}(t) dt$$

Resulting expression:

$$\langle \nabla_{\mathbf{b}} J, \tilde{\mathbf{b}} \rangle = \sum_{i=1}^I \int_0^T (\boldsymbol{\varphi}^i(t))^\top \tilde{\mathbf{b}}(t) dt + w_2 \int_0^T (\mathbf{b}(t))^\top \tilde{\mathbf{b}}(t) dt$$

Resulting gradient (w.r.t. the standard L^2 -innerproduct):

$$(\nabla_{\mathbf{b}} J)(t) = \sum_{i=1}^I \boldsymbol{\varphi}^i(t) + w_2 \mathbf{b}(t)$$

The gradient w.r.t. $\mathbf{V}(t)$

$$\begin{aligned} \langle \nabla_{\mathbf{V}} J, \hat{\mathbf{V}} \rangle_F &= \lim_{h \rightarrow 0} \frac{J(\mathbf{V} + h\hat{\mathbf{V}}, \mathbf{b}) - J(\mathbf{V}, \mathbf{b})}{h} = \sum_{i=1}^I (\mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i)^\top \hat{\mathbf{x}}^i(T) + \\ &\quad w_1 \sum_{i=1}^I \int_0^T (\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i)^\top \hat{\mathbf{x}}^i(t) \, dt + w_2 \int_0^T \langle \mathbf{V}(t), \hat{\mathbf{V}}(t) \rangle_F \, dt. \end{aligned}$$

We can now verify similarly as before that

$$(\mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i)^\top \hat{\mathbf{x}}(T) = -w_1 \int_0^T (\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i)^\top \hat{\mathbf{x}}^i(t) \, dt + \int_0^T (\boldsymbol{\varphi}^i(t))^\top \hat{\mathbf{V}}(t) \sigma(\mathbf{x}^i(t)) \, dt$$

In a similar way, we can show that the gradient w.r.t. $\mathbf{V}(t)$ (w.r.t. the Frobenius inner product) is

$$(\nabla_{\mathbf{V}} J)(t) = \sum_{i=1}^I \sigma(\mathbf{x}^i(t)) (\boldsymbol{\varphi}^i(t))^\top + w_2 \mathbf{V}(t).$$

An algorithm for the computation of the gradients

Computation of the gradients $\nabla_{\mathbf{V}} J(\mathbf{V}, \mathbf{b})$ and $\nabla_{\mathbf{b}} J(\mathbf{V}, \mathbf{b})$ (gradient in the point (\mathbf{V}, \mathbf{b}))

- Compute for $i = 1, 2, 3, \dots, I$ the solutions of

$$\mathbf{x}^i(0) = \mathbf{x}_{\text{in}}^i, \quad \dot{\mathbf{x}}^i(t) = \mathbf{V}(t)\sigma(\mathbf{x}^i(t)) + \mathbf{b}(t).$$

- Compute for $i = 1, 2, 3, \dots, I$ the solutions of

$$\varphi^i(T) = \mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i, \quad -\dot{\varphi}^i(t) = \left(\mathbf{V}(t) \text{diag} \left(\frac{d\sigma}{dx}(\mathbf{x}^i(t)) \right) \right)^\top \varphi^i(t) + w_1(\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i),$$

- The gradients are now given by

$$(\nabla_{\mathbf{V}} J)(t) = \sum_{i=1}^I \sigma(\mathbf{x}^i(t))(\varphi^i(t))^\top + w_2 \mathbf{V}(t), \quad (\nabla_{\mathbf{b}} J)(t) = \sum_{i=1}^I \varphi^i(t) + w_2 \mathbf{b}(t)$$

An algorithm for the computation of the gradients

Computation of the gradients $\nabla_{\mathbf{V}} J(\mathbf{V}, \mathbf{b})$ and $\nabla_{\mathbf{b}} J(\mathbf{V}, \mathbf{b})$ (gradient in the point (\mathbf{V}, \mathbf{b}))

- Compute for $i = 1, 2, 3, \dots, I$ the solutions of

$$\mathbf{x}^i(0) = \mathbf{x}_{\text{in}}^i, \quad \dot{\mathbf{x}}^i(t) = \mathbf{V}(t)\sigma(\mathbf{x}^i(t)) + \mathbf{b}(t).$$

- Compute for $i = 1, 2, 3, \dots, I$ the solutions of

$$\varphi^i(T) = \mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i, \quad -\dot{\varphi}^i(t) = \left(\mathbf{V}(t) \text{diag} \left(\frac{d\sigma}{dx}(\mathbf{x}^i(t)) \right) \right)^\top \varphi^i(t) + w_1(\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i),$$

- The gradients are now given by

$$(\nabla_{\mathbf{V}} J)(t) = \sum_{i=1}^I \sigma(\mathbf{x}^i(t))(\varphi^i(t))^\top + w_2 \mathbf{V}(t), \quad (\nabla_{\mathbf{b}} J)(t) = \sum_{i=1}^I \varphi^i(t) + w_2 \mathbf{b}(t)$$

Remark: when I is large, we need to solve many ODEs at each iteration. As we also need many iterations (e.g. 10,000) this can lead to a huge computational cost.

Stochastic gradient descent methods reduce this cost by considering a randomly selected subset of indices $i \in \{1, 2, \dots, I\}$ at each time step.

We will not go into this further in this lecture.

2.C Training of deep residual neural networks



Cost functional and dynamics

$$\min_{\mathbf{V}, \mathbf{b}} J(\mathbf{V}, \mathbf{b}) = \frac{1}{2} \sum_{i=1}^I |\mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i|^2 +$$

$$\frac{w_1}{2} \sum_{i=1}^I \int_0^T |\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i|^2 dt + \frac{w_2}{2} \int_0^T (\|\mathbf{V}(t)\|_F^2 + |\mathbf{b}(t)|^2) dt.$$

subject to the dynamics (for $i = 1, 2, 3, \dots, I$)

$$\mathbf{x}^i(0) = \mathbf{x}_{\text{in}}^i, \quad \dot{\mathbf{x}}^i(t) = \mathbf{V}(t)\sigma(\mathbf{x}^i(t)) + \mathbf{b}(t).$$

We consider a uniform grid $t_k = (k - 1)\Delta t$ ($k = 1, 2, \dots, N_T$), so $\Delta t = T/(N_T - 1)$.

We denote $\mathbf{x}_k^i \approx \mathbf{x}^i(t_k)$ with $k = 1, 2, \dots, N_T$,

$\mathbf{V}_k = \mathbf{V}(t_k)$ and $\mathbf{b}_k = \mathbf{b}(t_k)$ with $k = 1, 2, \dots, N_T - 1$.

We discretize with forward Euler:

$$\min_{\mathbf{V}_k, \mathbf{b}_k} J(\mathbf{V}, \mathbf{b}) = \frac{1}{2} \sum_{i=1}^I |\mathbf{x}_{N_T}^i - \mathbf{y}_{\text{out}}^i|^2 +$$

$$\frac{w_1 \Delta t}{2} \sum_{i=1}^I \sum_{k=2}^{N_T-1} |\mathbf{x}_k^i - \mathbf{y}_{\text{out}}^i|^2 + \frac{w_2 \Delta t}{2} \sum_{k=1}^{N_T-1} (\|\mathbf{V}_k\|_F^2 + |\mathbf{b}_k|^2).$$

$$\mathbf{x}_1^i = \mathbf{x}_{\text{in}}^i, \quad \mathbf{x}_{k+1}^i = \mathbf{x}_k^i + \Delta t(\mathbf{V}_k \sigma(\mathbf{x}_k^i) + \mathbf{b}_k).$$

Note: Forward Euler gives us precisely the structure of a ResNet.

Adjoint state

In the continuous time setting, we could compute the gradient from the adjoint state:

$$\varphi^i(T) = \mathbf{x}^i(T) - \mathbf{y}_{\text{out}}^i, \quad -\dot{\varphi}^i(t) = \left(\mathbf{V}(t) \text{diag} \left(\frac{d\sigma}{dx}(\mathbf{x}^i(t)) \right) \right)^\top \varphi^i(t) + w_1(\mathbf{x}^i(t) - \mathbf{y}_{\text{out}}^i),$$

Adjoint variables are $\varphi_k \approx \varphi(t_k)$, $k = 1, 2, 3, \dots, N_T - 1$.

Compute the adjoint variables starting from

$$\varphi_{N_T-1}^i = \mathbf{x}_{N_T}^i - \mathbf{y}_{\text{out}}^i$$

and then backward in time according to

$$\varphi_{k-1}^i = \varphi_k^i + \Delta t \left(\mathbf{V}_k \text{diag} \left(\frac{d\sigma}{dx}(\mathbf{x}_k^i) \right) \right)^\top \varphi_k^i + \Delta t w_1(\mathbf{x}_k^i - \mathbf{y}_{\text{out}}^i).$$

Gradient computation

In the continuous time setting:

$$(\nabla_{\mathbf{V}} J)(t) = \sum_{i=1}^I \sigma(\mathbf{x}^i(t)) (\boldsymbol{\varphi}^i(t))^\top + w_2 \mathbf{V}(t),$$

$$(\nabla_{\mathbf{b}} J)(t) = \sum_{i=1}^I \boldsymbol{\varphi}^i(t) + w_2 \mathbf{b}(t)$$

After discretization:

$$(\nabla_{\mathbf{V}} J)_k = \Delta t \sum_{i=1}^I \boldsymbol{\varphi}_k^i \sigma(\mathbf{x}_k^i)^\top + w_2 \Delta t \mathbf{V}_k, \quad k = 1, 2, \dots, N_T - 1,$$

$$(\nabla_{\mathbf{b}} J)_k = \Delta t \sum_{i=1}^I \boldsymbol{\varphi}_k^i + w_2 \Delta t \mathbf{b}_k, \quad k = 1, 2, \dots, N_T - 1.$$

Optimization algorithm

We can now just use a basic gradient descent algorithm to minimize J .
In every iteration we thus use the updates


$$\mathbf{V}_k^{j+1} = \mathbf{V}_k^j - \beta_j (\nabla_{\mathbf{V}} J)_k, \quad \mathbf{b}_k^{j+1} = \mathbf{b}_k^j - \beta_j (\nabla_{\mathbf{b}} J)_k.$$

The stepsize γ_j is also called the learning rate.

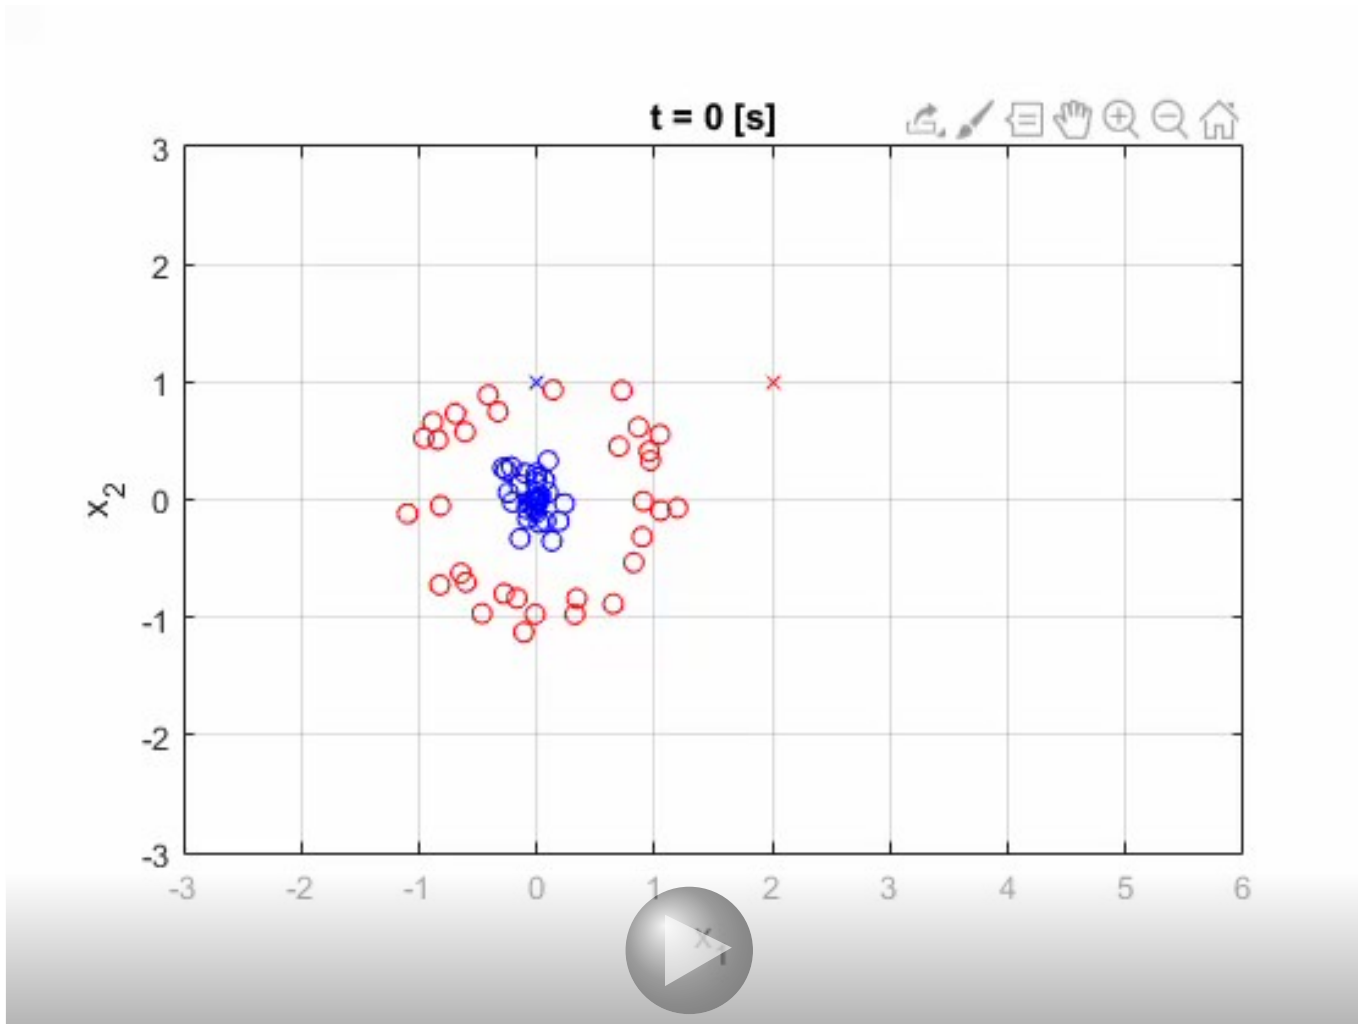
But the problem is now very much nonconvex:

- ▶ We cannot guarantee the uniqueness of the (global) minimizer.
- ▶ We do not know whether the algorithm converges to a global minimizer.
- ▶ The convergence rate is generally slow.
We need many iterations (e.g. 10,000) to obtain good results.

Example: 100 iterations of gradient descent



Example: 1000 iterations of gradient descent



Example: 10,000 iterations of gradient descent

